

A Scalable Method for Multiagent Constraint Optimization

Adrian Petcu and Boi Faltings
{adrian.petcu, boi.faltings}@epfl.ch
<http://liawww.epfl.ch/>

Artificial Intelligence Laboratory
Ecole Polytechnique Fédérale de Lausanne (EPFL)
IN (Ecublens), CH-1015 Lausanne, Switzerland

Abstract

We present in this paper a new, complete method for distributed constraint optimization, based on *dynamic programming*. It is a utility propagation method, inspired by the *sum-product algorithm*, which is correct only for tree-shaped constraint networks. In this paper, we show how to extend that algorithm to arbitrary topologies using a pseudotree arrangement of the problem graph. Our algorithm requires a *linear number of messages*, whose maximal size depends on the *induced width* along the particular pseudotree chosen.

We compare our algorithm with backtracking algorithms, and present experimental results. For some problem types we report orders of magnitude fewer messages, and the ability to deal with arbitrarily large problems. Our algorithm is formulated for optimization problems, but can be easily applied to satisfaction problems as well.

1 Introduction

Distributed Constraint Satisfaction (DisCSP) was first studied by Yokoo [Yokoo *et al.*, 1992] and has recently attracted increasing interest. In distributed constraint satisfaction each variable and constraint is owned by an agent. Systematic search algorithms for solving DisCSP are generally derived from depth-first search algorithms based on some form of backtracking [Silaghi *et al.*, 2000; Yokoo *et al.*, 1998; Meisels and Zivan, 2003; Hamadi *et al.*, 1998].

Recently, the paradigm of asynchronous distributed search has been extended to constraint optimization by integrating a bound propagation mechanism (ADOPT - [Modi *et al.*, 2003]).

In general, optimization problems are much harder to solve than DisCSP ones, as the goal is not just to find *any* solution, but the *best* one, thus requiring more exploration of the search space. The common goal of all distributed algorithms is to minimize the number of messages required to find a solution.

Backtracking algorithms are very popular in centralized systems because they require very little memory. In a distributed implementation, however, they may not be the best basis since in backtrack search, control shifts rapidly between

different variables. Every state change in a distributed backtrack algorithm requires at least one message; in the worst case, even in a parallel algorithm there will be exponentially many state changes [Kasif, 1986], thus resulting in exponentially many messages. So far, this has been a serious drawback for the application of distributed algorithms in the real world, especially for optimization problems (also noted in [Maheswaran *et al.*, 2004]).

This leads us to believe that other search paradigms, in particular those based on dynamic programming, may be more appropriate for DisCSP. For example, an algorithm that incrementally computes the set of all partial solutions for all previous variables according to a certain order would only use a linear number of messages. However, the messages could grow exponentially in size, and the algorithm would not have any parallelism.

Recently, the sum-product algorithm [Kschischang *et al.*, 2001] has been proposed for certain constraint satisfaction problems, for example decoding. It is an acceptable compromise as it combines a dynamic-programming style exploration of a search space with a fixed message size, and can easily be implemented in a distributed fashion. However, it is correct only for tree-shaped constraint networks.

In this paper, we show how to extend the algorithm to arbitrary topologies using a pseudotree arrangement of the problem graph, and report our experimental results. The algorithm is formulated for optimization problems, but can be easily applied to satisfaction problems by having relations with utility either 0 (for allowed tuples) or negative values (for disallowed tuples). Utility maximization produces a solution if there is an assignment with utility 0.

The rest of this paper is structured as follows: Section 2 presents the definitions and the notation we use, Section 3 presents an optimization procedure for trees, Section 4 the optimization for graphs, Section 5 proves the complexity to be equal to the induced width, Section 6 compares theoretically our algorithm with other approaches, Section 7 presents experimental results, and we conclude in Section 8.

2 Definitions & notation

A discrete *multiagent constraint optimization problem* (MCOP) is a tuple $\langle \mathcal{X}, \mathcal{D}, \mathcal{R} \rangle$ such that:

- $\mathcal{X} = \{X_1, \dots, X_m\}$ is the set of variables/agents;

- $\mathcal{D} = \{d_1, \dots, d_m\}$ is a set of domains of the variables, each given as a finite set of possible values.
- $\mathcal{R} = \{r_1, \dots, r_p\}$ is a set of relations, where a relation r_i is a function $d_{i1} \times \dots \times d_{ik} \rightarrow \mathbb{R}^+$ which denotes how much utility is assigned to each possible combination of values of the involved variables.

In this paper we deal with unary and binary relations, being well-known that higher arity relations can also be expressed in these terms with little modifications. In a MCOP, any value combination is allowed; the goal is to find an assignment \mathcal{X}^* for the variables X_i that maximizes the sum of utilities.

For a node X_k , we define $R_k(X_j) =$ the relation(s) between X_k and its neighbor X_j .

3 Distributed constraint optimization for tree-structured networks

For tree-structured networks, polynomial-time complete optimization methods have been developed (e.g. the sum-product algorithm [Kschischang *et al.*, 2001] and the *DTREE* algorithm from [Petcu and Faltings, 2004]).

In *DTREE*, the agents send *UTIL* messages (utility vectors) to their parents. A child X_l of node X_k would send X_k a vector of the optimal utilities $u_{X_l}^*(v_k^j)$ that can be achieved by the subtree rooted at X_l plus $R_l(X_k)$, and are compatible with each value v_k^j of X_k (such a vector has $|dom(X_k)|$ values).

For the leaf nodes it is immediate to compute these valuations by just inspecting the constraints they have with their single neighbors, so they initiate the process. Then each node X_i relays these messages according to the following process:

- Wait for *UTIL* messages from all children. Since all of the respective subtrees are disjoint, by summing them up, X_i computes how much utility each of its values gives for the whole subtree rooted at itself. This, together with the relation(s) between X_i and its parent X_j , enables X_i to compute exactly how much utility can be achieved by the entire subtree rooted at X_i , taking into account compatibility with each of X_j 's values. Thus, X_i can send to X_j its *UTIL* message. X_i also stores its optimal values corresponding to each value of X_j .
- If root node, X_i can compute the optimal overall utility corresponding to each one of its values (based on all the incoming *UTIL* messages), pick the optimal one, and send a *VALUE* message to its children, informing them about its decision.

Upon receipt of the *VALUE* message from its parent, each node is able to pick the optimal value for itself (as the previously stored optimal value corresponding to the value its parent has chosen), and pass it on to its children. At this point, the algorithm is finished for X_i .

The correctness of this algorithm was shown in the original paper, as well as the fact that it requires a linear number of messages, and linear memory.




Figure 1: Example of a pseudotree arrangement.

4 Distributed constraint optimization for general networks

To apply a *DTREE*-like algorithm to a cyclic graph, we first need to arrange the graph as a pseudotree (it is known that this arrangement is possible for any graph).

4.1 Pseudotrees

Definition 1 A pseudo-tree arrangement of a graph G is a rooted tree with the same vertices as G and the property that adjacent vertices from the original graph fall in the same branch of the tree (e.g. X_0 and X_{11} in Figure 1).

Pseudotrees have already been investigated as a means to boost search ([Freuder, 1985; Freuder and Quinn, 1985; Dechter, 2003; Schiex, 1999]). The main idea with their use in search, is that due to the relative independence of nodes lying in different branches of the pseudotree, it is possible to perform search in parallel on these independent branches.

Figure 1 shows an example of a pseudotree that we shall refer to in the rest of this paper. It consists of *tree edges*, shown as solid lines, and *back edges*, shown as dashed lines, that are not part of the spanning tree (e.g. $8-1$, $12-2$, $4-0$). We call a path in the graph that is entirely made of tree edges, a *tree-path*. A *tree-path associated with a back-edge* is the tree-path connecting the two nodes involved in the back-edge (please note that since our arrangement is a pseudotree, such a tree path is always included in a branch of the tree). For each back-edge, the higher node involved in that back-edge is called the *back-edge handler* (e.g. the dark nodes 0, 1 and 2). We also define:

- $P(X)$ - the *parent* of a node X : the single node higher in the hierarchy of the pseudotree that is connected to the node X directly through a tree edge (e.g. $P(X_4) = X_1$)
- $C(X)$ - the *children* of a node X : the set of nodes lower in the pseudotree that are connected to the node X directly through tree edges (e.g. $C(X_1) = \{X_3, X_4\}$)
- $PP(X)$ - the *pseudo-parents* of a node X : the set of nodes higher in the pseudotree that are connected to the node X directly through back-edges ($PP(X_8) = \{X_1\}$)
- $PC(X)$ - the *pseudo-children* of a node X : the set of nodes lower in the hierarchy of the pseudotree that are connected to the node X directly through back-edges (e.g. $PC(X_0) = \{X_4, X_{11}\}$)

As it is already known, a DFS (depth-first search) tree is also a pseudotree (although the inverse does not always hold). So, a DFS tree obtained from the DFS traversal of the graph starting from one of the nodes (chosen through a distributed leader election algorithm) will do just fine. Due to the lack of space we do not present here a procedure for the creation of a DFS tree, and refer the reader to techniques like [Barbosa, 1996; Hamadi *et al.*, 1998].

4.2 The DPOP algorithm

Our algorithm has 3 phases. First, the agents establish the pseudotree structure (see section 4.1) to be used in the following two phases. The next two phases are the *UTIL* and *VALUE* propagations, which are similar to the ones from *DTREE* - section 3. Please refer to Algorithm 1 for a formal description of the algorithm, and to the rest of this section for a detailed description of the *UTIL* and *VALUE* phases.

UTIL propagation

As in *DTREE*, the *UTIL* propagation starts from the leaves of the pseudotree and propagates up the pseudotree, only through the tree edges. It is easy for an agent to identify whether it is a leaf in the pseudotree or not: it must have a single tree edge (e.g. X_7 to X_{13} in Figure 1).

In a tree network, a *UTIL* message sent by a node to its parent is dependent only on the subtree rooted at the respective node (no links to other parts of the tree), and the constraint between the node and its parent. For example, consider the message ($X_6 \rightarrow X_2$). This message is clearly dependent only on the target variable X_2 , since there are no links between X_6 or X_{13} and any node above X_2 .

In a network with cycles (each back-edge in the pseudotree produces a cycle), a message sent from a node to its parent may also depend on variables above the parent. This happens when there is a back-edge connecting the sending node with such a variable. For example, consider the message ($X_8 \rightarrow X_3$) in Figure 1. We see that the utilities that the subtree rooted at X_8 can achieve are not dependent only on its parent X_3 (as for $X_6 \rightarrow X_2$). As X_8 is connected with X_1 through the backedge $X_8 \rightarrow X_1$, X_8 must take into account this dependency when sending its message to X_3 .

This is where the dynamic programming approach comes into play: X_8 will compute the optimal utilities its subtree can achieve for each value combination of the tuple $\langle X_3, X_1 \rangle$. It will then assemble a message as a hypercube with 2 dimensions (one for the target variable X_3 and one for the back-edge handler X_1), and send it to X_3 (see Table 1).

This is the key difference between *DTREE* and *DPOP*: messages travelling through the network in *DTREE* always have a single dimension (they are linear in the domain size of the target variable), whereas in *DPOP*, messages have multiple dimensions (one for the target variable, and another one for each context variable).

Combining messages - dimensionality increase

Let us consider this example: X_5 receives 2 messages from its children X_{11} and X_{12} ; the message from X_{11} has X_0 as context, and the one from X_{12} has X_2 as context. Both have one dimension for X_5 (target variable) and one dimension for their context variable (X_0 and X_2 respectively), therefore,

Algorithm 1: DPOP - Distributed pseudotree-optimization procedure for general networks.

```

1: DPOP( $\mathcal{X}, \mathcal{D}, \mathcal{R}$ )
   Each agent  $X_i$  executes:
2:
3: Phase 1: pseudotree creation
4: elect leader from all  $X_j \in \mathcal{X}$ 
5: elected leader initiates pseudotree creation
6: afterwards,  $X_i$  knows  $P(X_i)$ ,  $PP(X_i)$ ,  $C(X_i)$  and  $PC(X_i)$ 
7: Phase 2: UTIL message propagation
8: if  $|Children(X_i)| == 0$  (i.e.  $X_i$  is a leaf node) then
9:    $UTIL_{X_i}(P(X_i)) \leftarrow \text{Compute\_utils}(P(X_i), PP(X_i))$ 
10:  Send_message( $P(X_i)$ ,  $UTIL_{X_i}(P(X_i))$ )
11: activate UTIL_Message_handler()
12: Phase 3: VALUE message propagation
13: activate VALUE_Message_handler()
14: END ALGORITHM
15:
16: UTIL_Message_handler( $X_k, UTIL_{X_k}(X_i)$ )
17: store  $UTIL_{X_k}(X_i)$ 
18: if UTIL messages from all children arrived then
19:   if Parent( $X_i$ ) == null (that means  $X_i$  is the root) then
20:      $v_i^* \leftarrow \text{Choose\_optimal}(\text{null})$ 
21:     Send VALUE( $X_i, v_i^*$ ) to all  $C(X_i)$ 
22:   else
23:      $UTIL_{X_i}(P(X_i)) \leftarrow \text{Compute\_utils}(P(X_i), PP(X_i))$ 
24:     Send_message( $P(X_i)$ ,  $UTIL_{X_i}(P(X_i))$ )
25:   return
26:
27: VALUE_Message_handler( $VALUE_{P(X_i)}^{X_i}$ )
28: add all  $X_k \leftarrow v_k^* \in VALUE_{P(X_i)}^{X_i}$  to agent_view
29:  $X_i \leftarrow v_i^* = \text{Choose\_optimal}(\text{agent\_view})$ 
30: Send VALUE_{X_i}^{X_i} to all  $X_l \in C(X_i)$ 
31:
32: Choose_optimal(agent_view)
33:
34:    $v_i^* \leftarrow \text{argmax}_{v_i} \sum_{X_l \in C(X_i)} UTIL_{X_l}(v_i, \text{agent\_view})$ 
35:
36: Compute_utils( $P(X_i)$ ,  $PP(X_i)$ )
37: for all combinations of values of  $X_k \in PP(X_i)$  do
38:   let  $X_j$  be Parent( $X_i$ )
39:   similarly to DTREE, compute a vector  $UTIL_{X_i}(X_j)$ 
   of all  $\{Util_{X_i}(v_i^*(v_j), v_j) | v_j \in Dom(X_j)\}$ 
40: assemble a hypercube  $UTIL_{X_i}(X_j)$  out of all these
   vectors (totaling  $|PP(X_i)| + 1$  dimensions).
41: return  $UTIL_{X_i}(X_j)$ 

```

$X_8 \rightarrow X_3$	$X_3 = v_3^0$	$X_3 = v_3^1$...	$X_3 = v_3^{m-1}$
$X_1 = v_1^0$	$u_{X_8}^*(v_1^0)$	$u_{X_8}^*(v_1^1)$...	$u_{X_8}^*(v_1^0)$
...
$X_1 = v_1^{n-1}$	$u_{X_8}^*(v_1^{n-1})$	$u_{X_8}^*(v_1^{n-1})$...	$u_{X_8}^*(v_1^{n-1})$

Table 1: *UTIL* message sent from X_8 to X_3 , in Figure 1

their dimensionality is 2. X_5 needs to send out its message to its parent (X_2). X_5 considers all possible values of X_2 , and for each one of them, all combinations of values of the context variables (X_0 and X_2) and X_5 are considered; the values of X_5 are always chosen such that the optimal utilities for each tuple $\langle X_0 \times X_2 \times X_5 \rangle$ are achieved. Note that since X_2 is both a context variable and the target variable, the message collapses to 2 dimensions, not 3.

One can think of this process as the cross product of messages $X_{11} \rightarrow X_5$ and $X_{12} \rightarrow X_5$ resulting in a hypercube with dimensions X_0, X_2 and X_5 , followed by a projection on the X_5 axis, which retains the optimal utilities for the tuples $\langle X_0 \times X_2 \rangle$ (optimizing w.r.t. X_5 given X_0 and X_2).

Collapsing messages - dimensionality decrease

Whenever a multi-dimensional *UTIL* message reaches a target variable that occupies one dimension in the message (a back-edge handler), the target variable optimizes itself out of the context, and the outgoing message loses the respective dimension.

We can take the example of X_1 , which is initially present in the context of the message $X_8 \rightarrow X_3$: once the message arrives at X_1 , since X_1 does not have any more influence on the upper parts of the tree, X_1 can "optimize itself away" by simply choosing the best value for itself, for each value of its parent X_0 (the normal *DTREE* process). Thus, one can see that a back edge handler (X_1 in our case) appears as an extra dimension in the messages travelling from the lower end of the back edge (X_8) to itself, through the tree path associated with the back edge ($X_8 \rightarrow X_3 \rightarrow X_1$).

VALUE propagation

The *VALUE* phase is similar to *DTREE*. Now, in addition to its parent's value, the $VALUE_{P(X_j)}^{X_j}$ message a node X_j receives from its parent also contains the values of all the variables that were present in the context of X_j 's *UTIL* message for its parent. E.g.: X_0 sends X_2 $VALUE_0^2(X_0 \leftarrow v_0^*)$, then X_2 sends X_5 $VALUE_2^5(X_0 \leftarrow v_0^*, X_2 \leftarrow v_2^*)$, and X_5 sends X_{11} $VALUE_5^{11}(X_0 \leftarrow v_0^*, X_5 \leftarrow v_5^*)$.

5 Complexity analysis

By construction, the number of messages our algorithm produces is linear: there are $n - 1$ *UTIL* messages - one through each tree-edge (n is the number of nodes in the problem), and m linear size *VALUE* messages - one through each edge (m is the number of edges). The DFS construction also produces a linear number of messages (good algorithms require $2 \times m$ messages). Thus, the complexity of this algorithm lies in the size of the *UTIL* messages.

Theorem 1 *The largest UTIL message produced by Algorithm 1 is space-exponential in the width of the pseudotree induced by the DFS ordering used.*

PROOF. Dechter ([Dechter, 2003], chapter 4, pages 86-88) describes the *fill-up method* for obtaining the *induced width*. First, we build the *induced graph*: we take the DFS traversal of the pseudotree as an ordering of the graph and process the nodes recursively (bottom up) along this order. When a node is processed, all its parents are connected (if not already

connected). The *induced width* is the maximum number of parents of any node in the induced graph.

It is shown in [Dechter, 2003] that the width of a tree (no back-edges) is 1. Actually the back-edges are the ones that influence the width: a single backedge produces an induced width of 2. From the construction of the induced tree, it is easy to see that several backedges produce increases in the width only when their tree-paths overlap on at least one edge, and their respective handlers are different; otherwise their effects on the width do not combine. Thus, the width is given by the size of the maximal set of back-edges which have overlapping tree-paths and distinct handlers.

During the *UTIL* propagation, the message size varies; the largest message is the one with the most dimensions. We have seen that a dimension X_i is added to a message when a back-edge with X_i as a handler is first encountered in the propagation, and travels through the tree-path associated with the back-edge. It is then eliminated by projection when the message arrives at X_i . The maximal dimensionality is therefore given by the maximal number of overlaps of tree-paths associated with back-edges with distinct handlers.

We have thus shown that the maximal dimensionality is equal to the induced width. \square

Exponential size messages are not necessarily a problem in all setups (depending on the resources available and on the induced width - low width problems generate small messages!)

However, when the maximum message size is limited, one can serialize big messages by letting the back-edge handlers ask explicitly for valuations for each one of their values *sequentially*, so each message can have customizable size.

A workaround against exponential memory is possible by renouncing exactness, and propagating valuations for the best/worst value combinations (upper/lower bounds) instead of all combinations.

6 Comparison with other approaches

Schiex [Schiex, 1999] notes the fact that so far, pseudotree arrangements have been mainly used for *search* procedures (essentially backtrack-based search, or branch-and-bound for optimization). As good examples, see the Distributed Depth-first Branch and Bound (DDBB), Distributed Iterative Deepening (DID), ADOPT, Synchronous Branch and Bound (SBB) and Iterative Distributed Breakout (IDB). All these procedures have a worst case complexity exponential in the depth of the pseudotree arrangement (basically because all the variables on the longest branch from root to a leaf have to be instantiated sequentially, and all their value combinations tried out). It was shown in [Bayardo and Miranker, 1995] that there are ways to obtain shallow pseudotrees (within a logarithmic factor of the induced width), but these require intricate heuristics like the ones from [Freuder and Quinn, 1985; Maheswaran *et al.*, 2004], which have not yet been adapted to a distributed setting, as also noted by the authors of the second paper.

In contrast, our approach exhibits a worst case complexity exponential in the *width of the graph* induced by the pseudotree ordering. Arnborg shows in [Arnborg, 1985] that finding a min-width ordering of a graph is NP-hard; however,

the DFS traversal of the graph has the advantage that it produces a good approximation, and is easy to implement in a distributed context. This, coupled with the fact that the depth of the pseudotree is irrelevant to the complexity, means that our algorithm works well with a simple DFS ordering. To see this fundamental difference between the two approaches, consider a problem that is a ring with n nodes. A DFS ordering of such a graph would yield a pseudotree with height n , and one back edge, thus the induced width is 2. A backtracking algorithm is time exponential in n , whereas our algorithm is linear, with message size $O(|d|^2)$. Since the exponential complexity translates directly in the explosion of the number of messages exchanged, these backtracking-based algorithms have not yet been applied to large systems.

Furthermore, it was shown by Dechter in [Dechter and Fattah, 2001] that the induced width is always less than or at most equal with the pseudotree height; thus we can conclude theoretically that our algorithm will always do at least as well as a pseudotree backtrack-based algorithm on the same pseudotree ordering. However, it is only fair to say that our approach can generate very big messages in the worst case, so one has to find a proper tradeoff between the number and the size of the messages transmitted through the system.

7 Experimental evaluation

Usual performance metrics for distributed algorithms are the number of messages and the number of synchronous cycles required to find the optimal solution. Both are linear in our case. For the number of messages, see section 5; the number of synchronous cycles is two times the height of the pseudotree (one *UTIL* propagation, and one *VALUE* propagation). We also introduce the maximal message size as a metric.

7.1 Sensor networks

One of our experimental setups is the sensor grid testbed from [Bejar *et al.*, 2005]. Briefly, there is a set of targets in a sensor field, and the problem is to allocate 3 different sensors to each target. This is a NP-complete resource allocation problem.

In [Bejar *et al.*, 2005], random instances are solved by AWC (a complete algorithm for constraint *satisfaction*). The problems are relatively small (100 sensors and maximum 18 targets, beyond which the problems become intractable). Our initial experiments with this setup solve to optimality problems in a grid of 400 sensors, with up to 40 targets.

Another setup is the one from [Maheswaran *et al.*, 2004], where there are corridors composed of squares which indicate areas to be observed. Sensors are located at each vertex of a square; in order for a square to be "observed", all 4 sensors in its vertices need to be focused on the respective square. Depending on the topology of the grid, some sensors are shared between several squares, and they can observe only one of them at a time. The authors test 4 improved versions of ADOPT (current state of the art for MCOP) on 4 different scenarios, where the corridors have the shapes of capital letters L, Z, T and H. Their results and a comparison with DPOP are in Table 2. One can see the dramatic reduction of the number of messages required (in some cases orders of magnitude), even for these very small problem instances (16

Algo/Scenario	Test L	Test Z	Test T	Test H
MCN , No Pass	626.4	1111.64	1841.28	1898.04
MLSP, No Pass	597.88	663.32	477.56	679.36
MCN , Pass	95.67	101.90	94.93	258.07
MLSP , Pass	81.77	91.5	107.77	255.2
DPOP	30	30	18	30

Table 2: DPOP vs 4 ADOPT versions: number of messages in sensor allocation problems.

variables). The explanation is that our algorithm always produces a linear number of messages.

Regarding the size of the messages: these problems have graphs with very low induced widths (2), basically given by the intersections between corridors. Thus, our algorithm employs linear messages in most of the parts of the problems, and only in the intersections are created 2 messages with 2 dimensions (in this case with 64 values each).

This fact gives our algorithm the ability to solve arbitrarily large instances of this particular kind of real-world problems.

7.2 Meeting scheduling

We experimented with distributed meeting scheduling in an organization with a hierarchical structure (a tree with departments as nodes, and a set of agents working in each department). The CSP model is the PEAV model from [Maheswaran *et al.*, 2004]. Each agent has multiple variables: one for the start time of each meeting it participates in, with 8 timeslots as values. Mutual exclusion constraints are imposed on the variables of an agent, and equality constraints are imposed on the corresponding variables of all agents involved in the same meeting. Private, unary constraints placed by an agent on its own variables show how much it values each meeting/start time. Random meetings are generated, each with a certain utility for each agent. The objective is to find the schedule that maximizes the overall utility.

Table 3 shows how our algorithm scales up with the size of the problems. Notice that the total number of messages includes the *VALUE* messages (linear size), and that due to the fact that intra-agent subproblems are denser than the rest of the problem, high-dimensional messages are likely to be virtual, intra-agent messages (not actually transmitted over the network). To our knowledge, these are by far the largest optimization problems solved with a complete, distributed algorithm (200 agents, 101 meetings, 270 variables, 341 constraints). The largest reported previous experiment is [Maheswaran *et al.*, 2004], with 33 agents, 12 meetings, 47 variables, 123 constraints, solved using ADOPT.

8 Conclusions and future work

We presented in this paper a new complete method for distributed constraint optimization. This method is a utility-propagation method that extends tree propagation algorithms like the sum-product algorithm or *DTREE* to work on arbitrary topologies using a pseudotree structure. It requires a *linear number of messages*, the largest one being exponential in the induced width along the particular pseudotree cho-

Agents	30	40	70	100	200
Meetings	14	15	34	50	101
Variables	44	50	112	160	270
Constraints	52	60	156	214	341
Messages	95	109	267	373	610
Max message size	512	4096	32k	256k	256k
Cycles	30	32	70	86	96

Table 3: DPOP tests on meeting scheduling.

sen. This method reduces the complexity from dom^n (standard backtracking) to dom^w , where n =number of nodes in the problem, dom bounds the domain size and w =the induced width along the particular pseudotree chosen. For loose problems, $n \gg w$ holds and our method retains the advantage of a linear number of messages (in practice even orders of magnitude fewer messages than the other approaches), while preserving a small message size. In real world scenarios, sending a few larger messages is preferable to sending a lot of small messages because of the much lower overheads implied (differences can go up to orders of magnitude speedups). Our experiments show that our method is the first one to be able to handle effectively arbitrarily large instances of a number of practical problems while using a linear number of messages.

Finding the minimum width pseudotree is an NP-complete problem, so in our future work we will investigate heuristics for finding low width pseudotrees.

9 Acknowledgements

We would like to thank Rina Dechter and Radu Marinescu for insightful discussions, Jonathan Pearce/TEAMCORE for experimental data from sensor networks and meeting scheduling simulations, and Michael Schumacher for valuable comments on an early version of this paper. This work has been funded by the Swiss National Science Foundation under contract No. 200020-103421/1.

References

- [Arnborg, 1985] S. Arnborg. Efficient algorithms for combinatorial problems on graphs with bounded decomposability - a survey. *BIT*, 25(1):2–23, 1985.
- [Barbosa, 1996] Valmir Barbosa. *An Introduction to Distributed Algorithms*. The MIT Press, 1996.
- [Bayardo and Miranker, 1995] Roberto Bayardo and Daniel Miranker. On the space-time trade-off in solving constraint satisfaction problems. In *Proceedings of the 15th International Joint Conference on Artificial Intelligence, IJCAI-95*, Montreal, Canada, 1995.
- [Bejar et al., 2005] Ramon Bejar, Cesar Fernandez, Magda Valls, Carmel Domshlak, Carla Gomes, Bart Selman, and Bhaskar Krishnamachari. Sensor networks and distributed CSP: Communication, computation and complexity. *Artificial Intelligence*, 161(1-2):117–147, 2005.
- [Dechter and Fattah, 2001] Rina Dechter and Yousri El Fattah. Topological parameters for time-space tradeoff. *Artificial Intelligence*, 125(1-2):93–118, 2001.
- [Dechter, 2003] Rina Dechter. *Constraint Processing*. Morgan Kaufmann, 2003.
- [Freuder and Quinn, 1985] Eugene C. Freuder and Michael J. Quinn. Taking advantage of stable sets of variables in constraint satisfaction problems. In *Proceedings of the 9th International Joint Conference on Artificial Intelligence, IJCAI-85*, pages 1076–1078, Los Angeles, CA, 1985.
- [Freuder, 1985] Eugene C. Freuder. A sufficient condition for backtrack-bounded search. *Journal of the ACM*, 32(14):755–761, 1985.
- [Hamadi et al., 1998] Y. Hamadi, C. Bessière, and J. Quinqueton. Backtracking in distributed constraint networks. In *ECAI-98*, pages 219–223, 1998.
- [Kasif, 1986] Simon Kasif. On the parallel complexity of some constraint satisfaction problems. In *Proceedings of the National Conference on Artificial Intelligence, AAAI-86*, pages 349–353, Philadelphia, PA, 1986.
- [Kschischang et al., 2001] Frank R. Kschischang, Brendan Frey, and Hans Andrea Loeliger. Factor graphs and the sum-product algorithm. *IEEE Transactions On Information Theory*, 47(2):498–519, FEBRUARY 2001.
- [Maheswaran et al., 2004] Rajiv T. Maheswaran, Milind Tambe, Emma Bowring, Jonathan P. Pearce, and Pradeep Varakantham. Taking DCOP to the realworld: Efficient complete solutions for distributed multi-event scheduling. In *AAMAS-04*, 2004.
- [Meisels and Zivan, 2003] Amnon Meisels and Roie Zivan. Asynchronous forward-checking on DisCSPs. In *Proceedings of the Distributed Constraint Reasoning Workshop, IJCAI 2003, Acapulco, Mexico*, 2003.
- [Modi et al., 2003] P. J. Modi, W. M. Shen, and M. Tambe. An asynchronous complete method for distributed constraint optimization. In *Proc. AAMAS*, 2003.
- [Petcu and Faltings, 2004] Adrian Petcu and Boi Faltings. A distributed, complete method for multi-agent constraint optimization. In *CP 2004 - Fifth International Workshop on Distributed Constraint Reasoning (DCR2004) in Toronto, Canada*, September 2004.
- [Schiex, 1999] Thomas Schiex. A note on CSP graph parameters. Technical Report 03, INRA, 07 1999.
- [Silaghi et al., 2000] Marius-Calin Silaghi, Djamila Sam-Haroud, and Boi Faltings. Asynchronous search with aggregations. In *AAAI/IAAI*, pages 917–922, Austin, Texas, 2000.
- [Yokoo et al., 1992] Makoto Yokoo, Edmund H. Durfee, Toru Ishida, and Kazuhiro Kuwabara. Distributed constraint satisfaction for formalizing distributed problem solving. In *International Conference on Distributed Computing Systems*, pages 614–621, 1992.
- [Yokoo et al., 1998] Makoto Yokoo, Edmund H. Durfee, Toru Ishida, and Kazuhiro Kuwabara. The distributed constraint satisfaction problem - formalization and algorithms. *IEEE Transactions on Knowledge and Data Engineering*, 10(5):673–685, 1998.